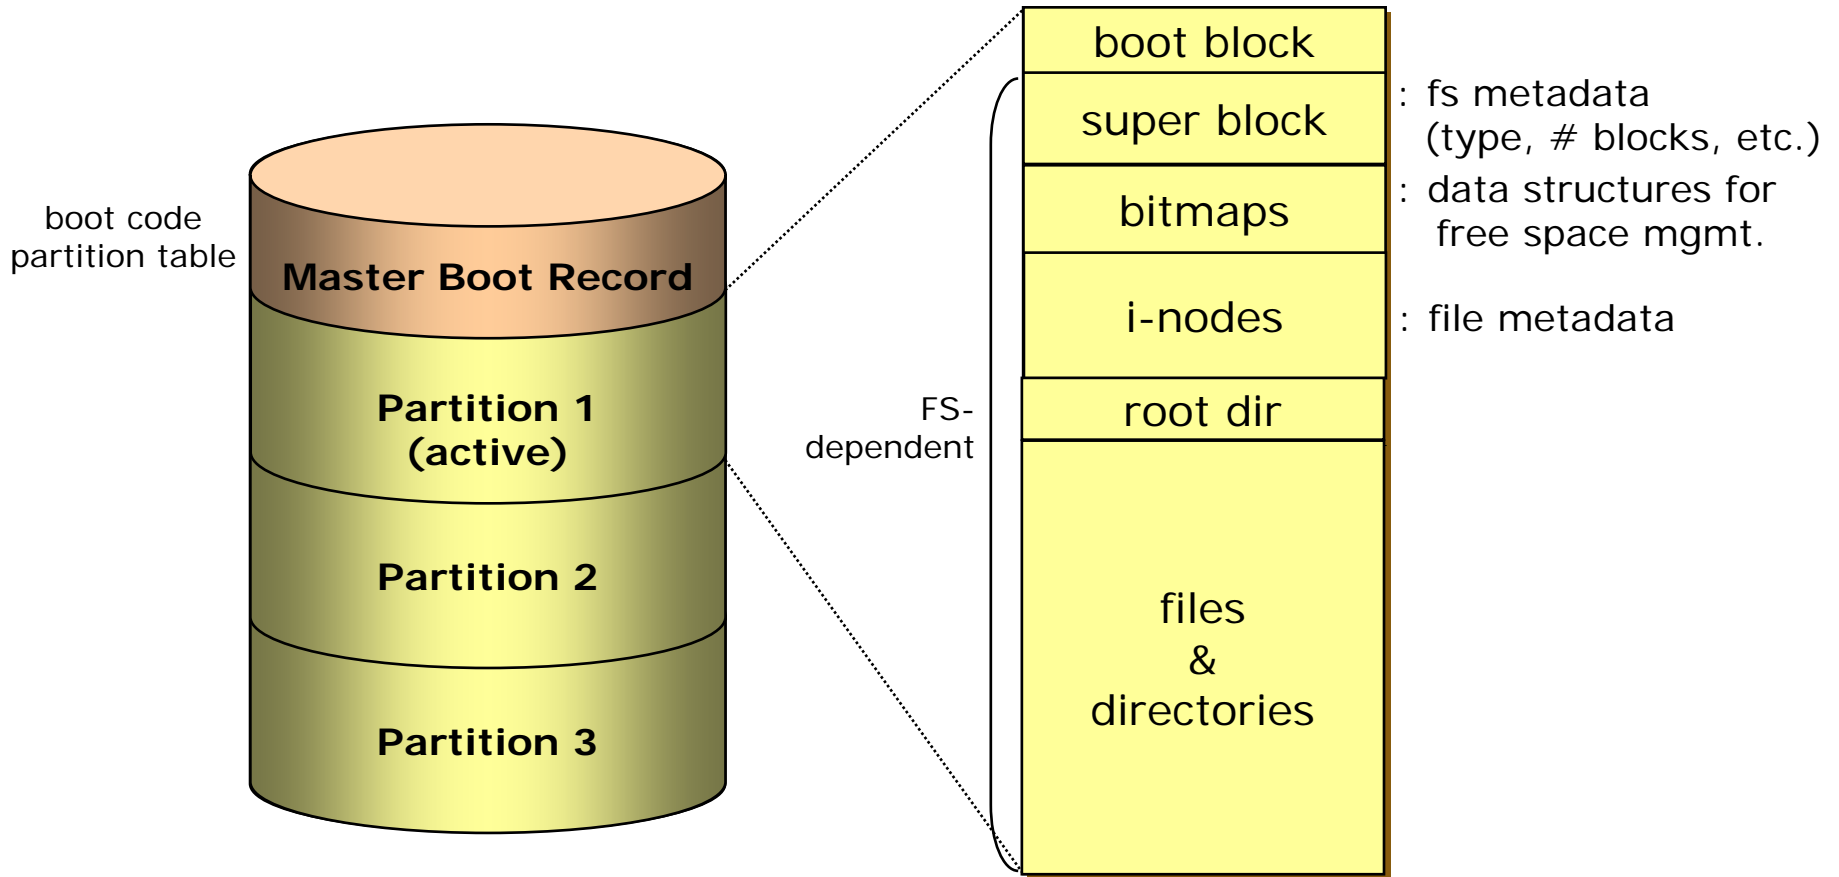


File System Implementation

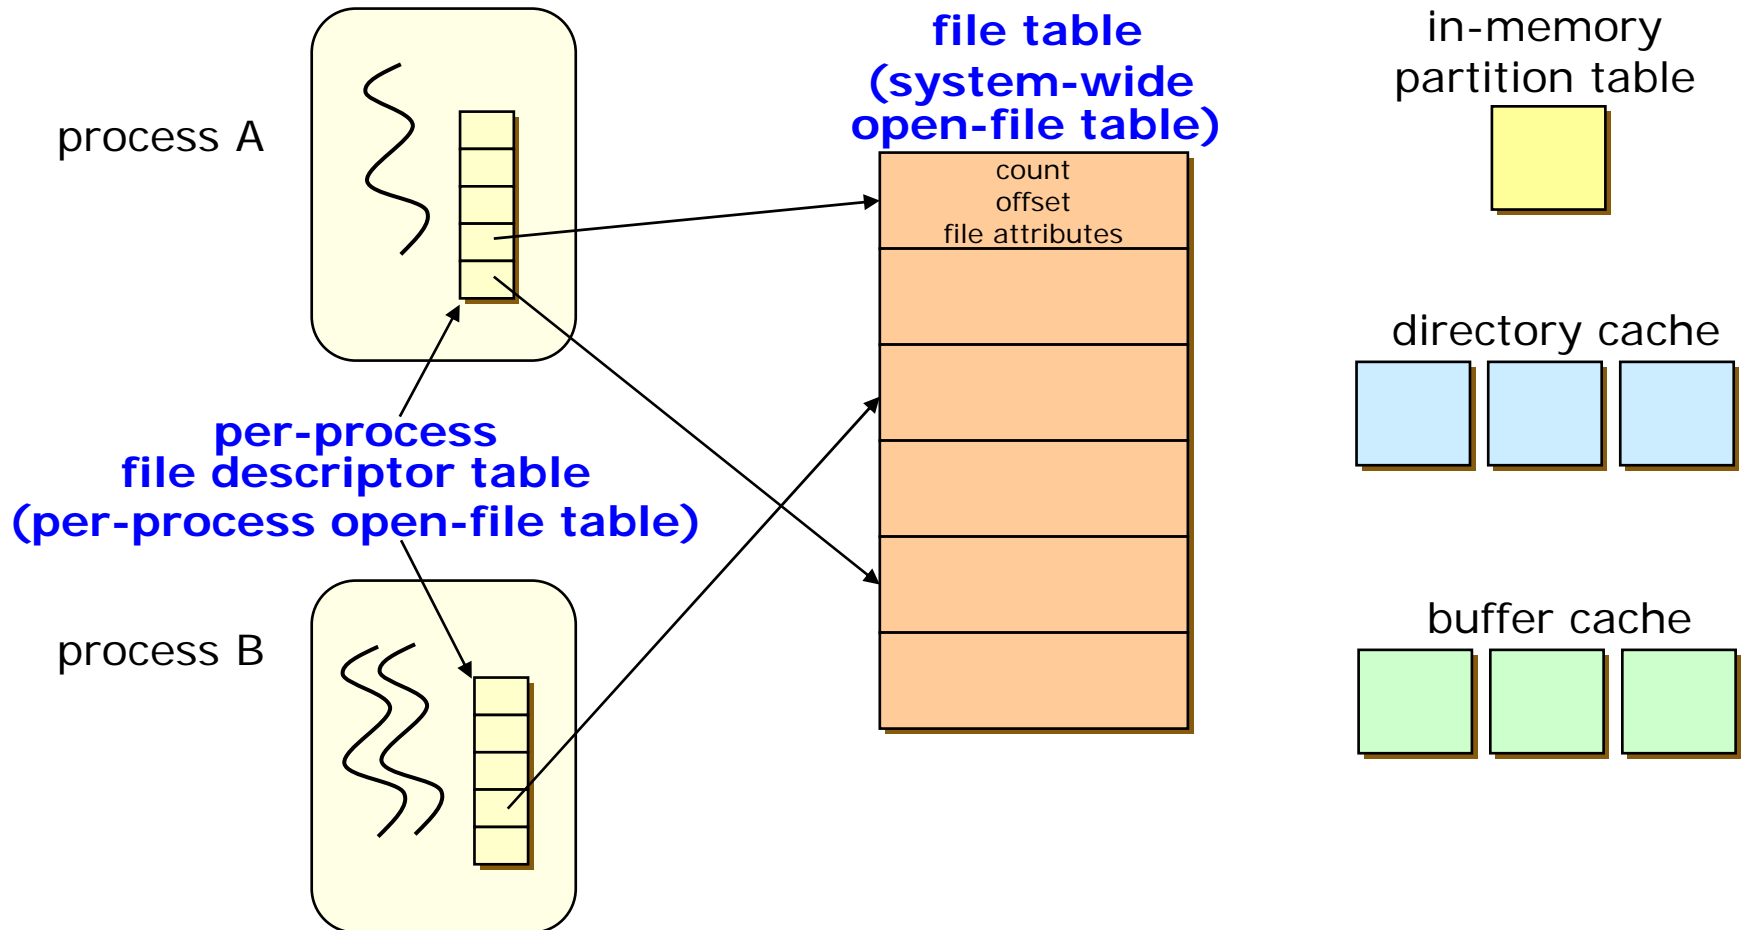
Overview

- User's view on file systems:
 - How files are named?
 - What operations are allowed on them?
 - What the directory tree looks like?
- Implementor's view on file systems:
 - How files and directories are stored?
 - How disk space is managed?
 - How to make everything work efficiently and reliably?

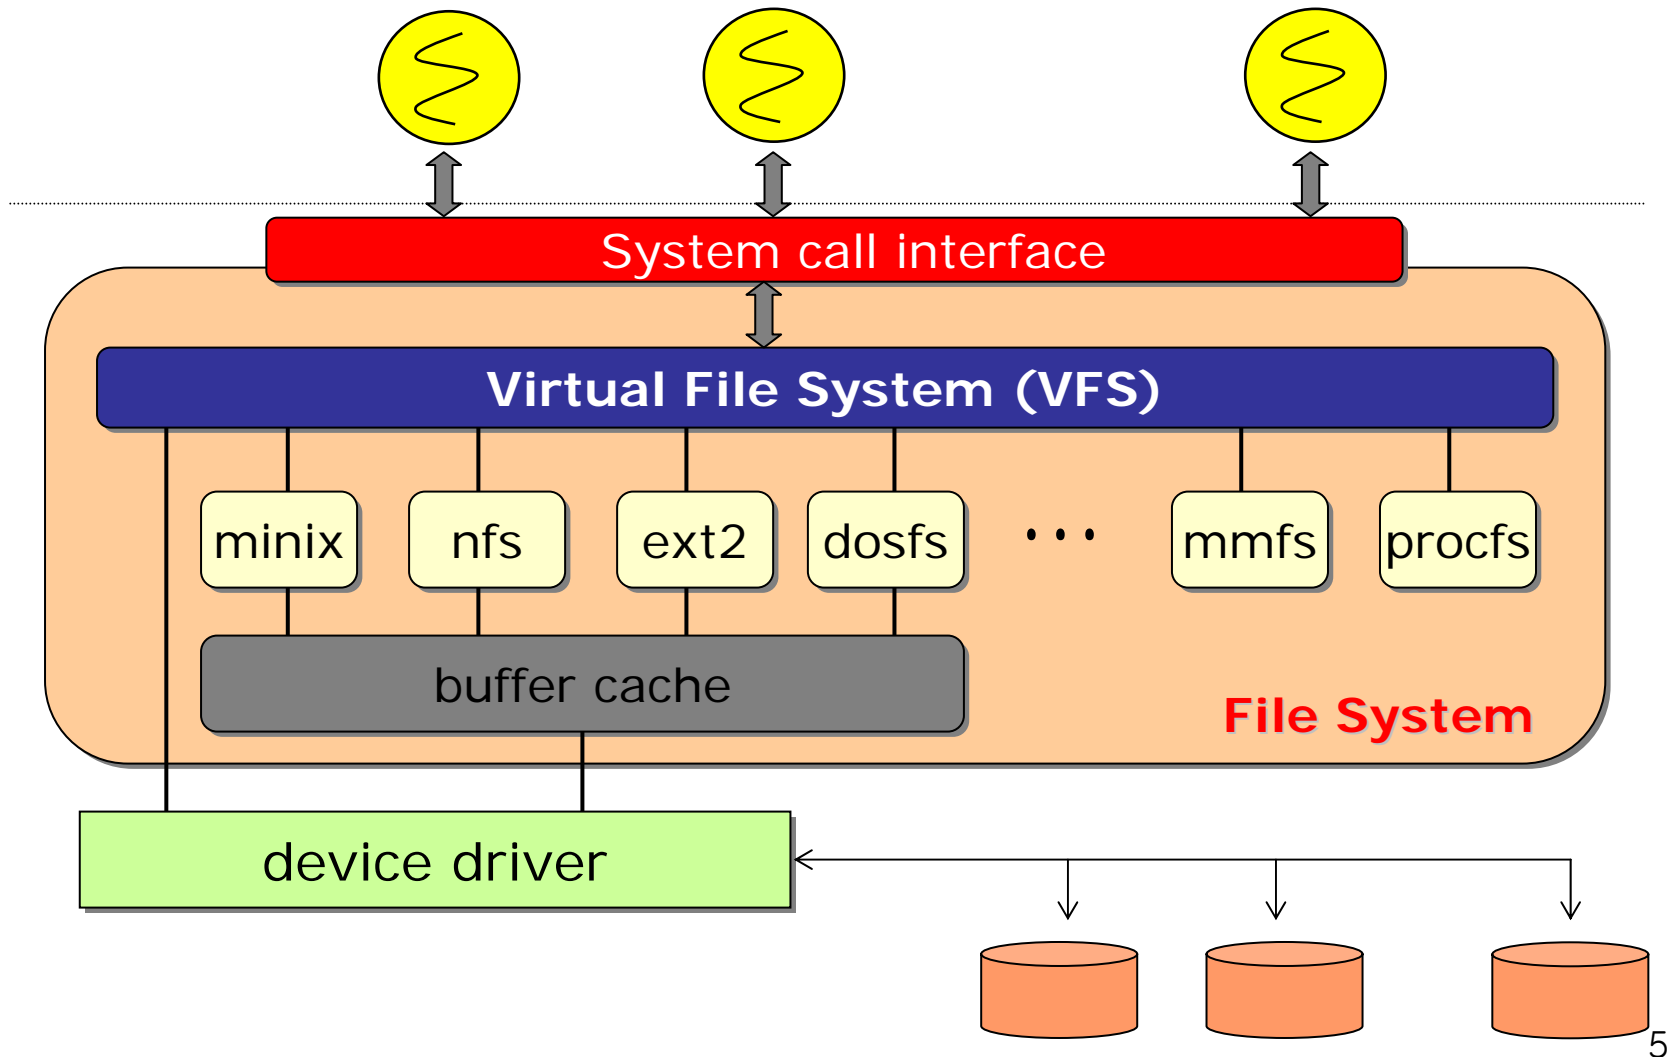
Disk Layout



In-memory Structures



File System Internals



VFS (1)

- Virtual File System
 - Manages kernel-level file abstractions in one format for all file systems.
 - Receives system call requests from user-level (e.g., open, write, stat, etc.)
 - Interacts with a specific file system based on mount point traversal.
 - Receives requests from other parts of the kernel, mostly from memory management.
 - Translates file descriptors to VFS data structures (such as vnode).

VFS (2)



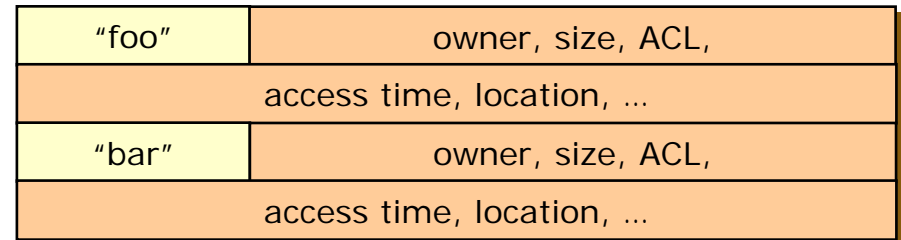
- Linux: VFS common file model
 - The superblock object
 - stores information concerning a mounted file system.
 - The inode object
 - stores general information about a specific file.
 - The file object
 - stores information about the interaction between an open file and a process.
 - The dentry object
 - stores information about the linking of a directory entry with the corresponding file.
 - In order to stick to the VFS common file model, in-kernel structures may be constructed on the fly.

Directory Implementation (1)

- Directory structure
 - Table (fixed length entries)
 - Linear list
 - Simple to program, but time-consuming.
 - Requires a linear search to find an entry.
 - Entries may be sorted to decrease the average search time and to produce a sorted directory listing easily (e.g., using B-tree).
 - Hash table
 - Decreases the directory search time.
 - A hash table is generally fixed size and the hash function depends on that size. (need mechanisms for collisions)
 - The number of files can be large:
 - (1) enlarge the hash table and remap.
 - (2) use a chained-overflow hash table.

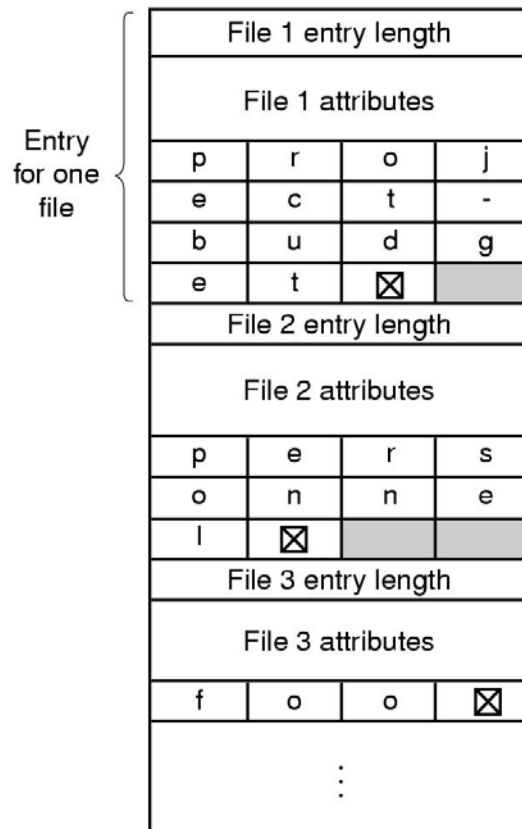
Directory Implementation (2)

- The location of metadata
 - In the directory entry
 - In the separate data structure (e.g., i-node)
 - A hybrid approach

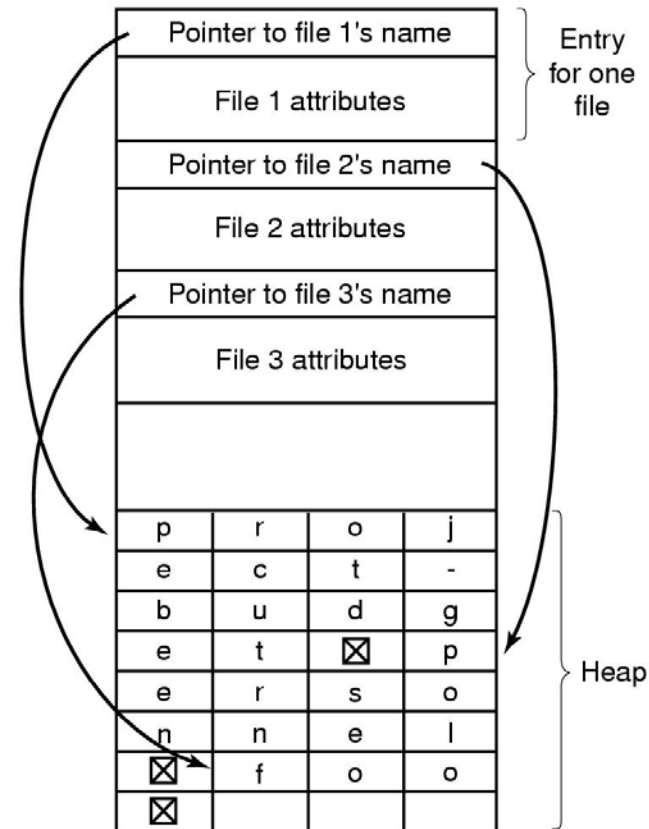


Directory Implementation (3)

- Supporting long file names




(a)



(b)

Allocation (1)

- Contiguous allocation
 - A file occupies a set of contiguous blocks on the disk.
 - Used by IBM VM/CMS



Allocation (2)


- Contiguous allocation (cont'd)
 - Advantages
 - The number of disk seeks is minimal.
 - Directory entries can be simple:
<file name, starting disk block, length, etc.>
 - Disadvantages
 - Requires a dynamic storage allocation: First / best fit.
 - External fragmentation: may require a compaction.
 - The file size is hard to predict and varying over time.
 - Feasible and widely used for CD-ROMS
 - All the file sizes are known in advance.
 - Files will never change during subsequent use.

Allocation (3)

- Modified contiguous allocation
 - A contiguous chunk of space is allocated initially.
 - When the amount is not large enough, another chunk of a contiguous space (an **extent**) is added.
 - Advantages
 - Still the directory entry can be simple.
<name, starting disk block, length, link to the extent>
 - Disadvantages
 - Internal fragmentation: if the extents are too large.
 - External fragmentation: if we allow varying-sized extents.
 - Used by Veritas File System (VxFS).

Allocation (4)

- Linked allocation
 - Each file is a linked list of disk blocks.



Allocation (5)

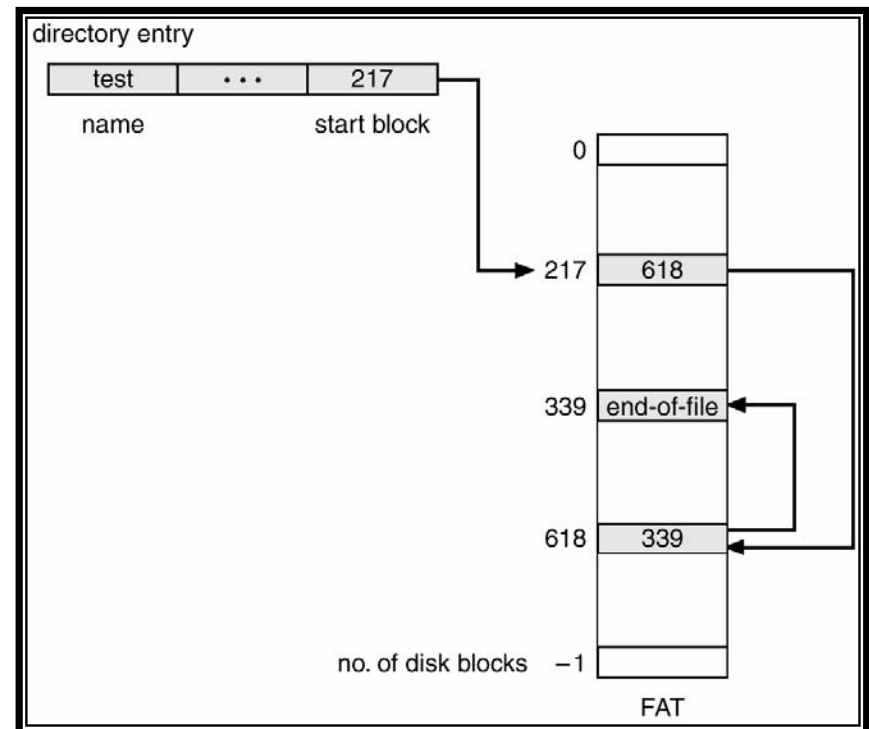
- Linked allocation (cont'd)
 - Advantages
 - Directory entries are simple:
<file name, starting block, ending block, etc.>
 - No external fragmentation: the disk blocks may be scattered anywhere on the disk.
 - A file can continue to grow as long as free blocks are available.
 - Disadvantages
 - It can be used only for sequentially accessed files.
 - Space overhead for maintaining pointers to the next disk block.
 - The amount of data storage in a block is no longer a power of two because the pointer takes up a few bytes.
 - Fragile: a pointer can be lost or damaged.

Allocation (6)

- Linked allocation using clusters
 - Collect blocks into multiples (clusters) and allocate the clusters to files.
 - e.g., 4 blocks / 1 cluster
 - Advantages
 - The logical-to-physical block mapping remains simple.
 - Improves disk throughput (fewer disk seeks)
 - Reduced space overhead for pointers.
 - Disadvantages
 - Internal fragmentation

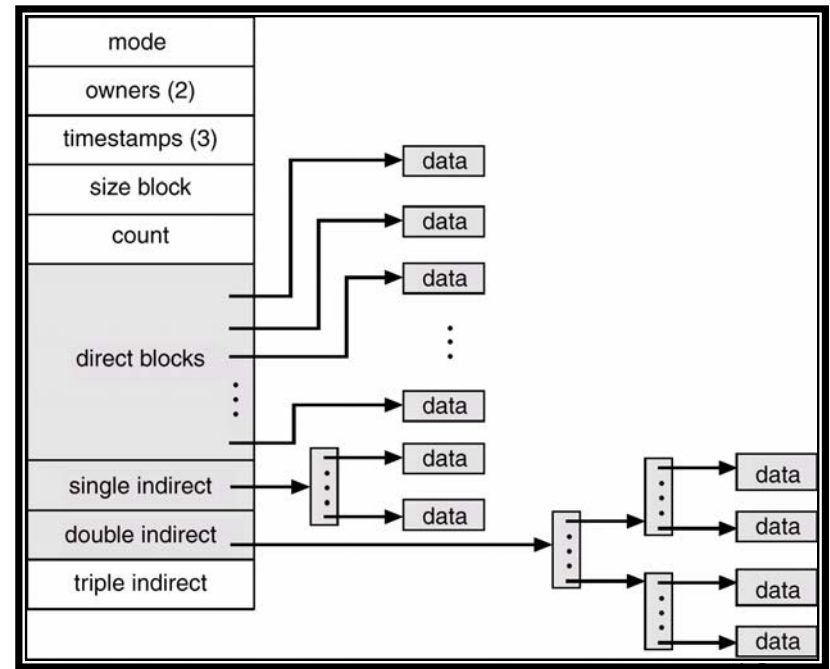

Allocation (7)

- Linked allocation using a FAT
 - A section of disk at the beginning of each partition is set aside to contain a file allocation table (FAT).
 - FAT should be cached to minimize disk seeks.
 - Space overhead can be substantial.
 - Random access time is improved.
 - Used by MS-DOS, OS/2
 - cf. FAT-16: 2GB limitation with 32KB block size



Allocation (8)

- Indexed allocation
 - Bring all the pointers together into one location (**index block** or **i-node**)
 - Each file has its own index block.



Allocation (9)


- Indexed allocation (cont'd)
 - Advantages
 - Supports direct access, without suffering from external fragmentation.
 - I-node need only be in memory when the corresponding file is open.
 - Disadvantages
 - Space overhead for indexes:
 - (1) Linked scheme: link several index blocks
 - (2) Multilevel index blocks
 - (3) Combined scheme: UNIX
 - 12 direct blocks, single indirect block, double indirect block, triple indirect block

Free Space Management (1)

- Bitmap or bit vector
 - Each block is represented by 1 bit.
 - 1 = free, 0 = allocated
 - Simple and efficient in finding the first free block.
 - May be accelerated by CPU's bit-manipulation instructions.
 - Inefficient unless the entire vector is kept in main memory.
 - Clustering reduces the size of bitmaps.

Free Space Management (2)

- Linked list
 - Link together all the free disk blocks, keeping a pointer to the first free blocks.
 - To traverse the list, we must read each block, but it's not a frequent action.
 - The FAT method incorporates free-block accounting into the allocation data structure.



Free Space Management (3)

- Grouping
 - Store the addresses of n free blocks in the first free block.
 - The addresses of a large number of free blocks can be found quickly.
- Counting
 - Keep the address of the free block and the number of free contiguous blocks.
 - The length of the list becomes shorter and the count is generally greater than 1.
 - Several contiguous blocks may be allocated or freed simultaneously.

Reliability (1)

- File system consistency
 - File system can be left in an inconsistent state if cached blocks are not written out due to the system crash.
 - It is especially critical if some of those blocks are i-node blocks, directory blocks, or blocks containing the free list.
 - Most systems have a utility program that checks file system consistency
 - Windows: scandisk
 - UNIX: fsck

Reliability (2)

- fsck: checking blocks
 - Reads all the i-nodes and mark used blocks.
 - Examines the free list and mark free blocks.

Consistent

	0	1	2	3	4	5	6	7
Blocks in use	1	1	0	1	0	1	1	1
Free blocks	0	0	1	0	1	0	0	0

Missing block

-- add it to the free list

	0	1	2	3	4	5	6	7
Blocks in use	1	1	0	1	0	1	1	1
Free blocks	0	0	0	0	1	0	0	0

Duplicated free block

-- rebuild the free list

	0	1	2	3	4	5	6	7
Blocks in use	1	1	0	1	0	1	1	1
Free blocks	0	0	1	0	2	0	0	0

Duplicated data block

-- allocate a new block and copy



	0	1	2	3	4	5	6	7
Blocks in use	1	1	0	1	0	2	1	1
Free blocks	0	0	1	0	1	0	0	0

Reliability (3)

- fsck: checking directories
 - Recursively descends the tree from the root directory, counting the number of links for each file.
 - Compare these numbers with the link counts stored in the i-nodes.
 - Force the link count in the i-node to the actual number of directory entries.

i-node count

1	1
5	2
12	4
...	...




Reliability (4)

- Journaling file systems
 - Fsck'ing takes a long time, which makes the file system restart slow in the event of system crash.
 - Record a log, or journal, of changes made to files and directories to a separate location. (preferably a separate disk).
 - If a crash occurs, the journal can be used to undo any partially completed tasks that would leave the file system in an inconsistent state.
 - IBM JFS for AIX, Linux
Veritas VxFS for Solaris, HP-UX, Unixware, etc.
SGI XFS for IRIX, Linux
Reiserfs, ext3 for Linux

Performance (1)

- Block size
 - Disk block size vs. file system block size
 - The median file size in UNIX is about 1KB.

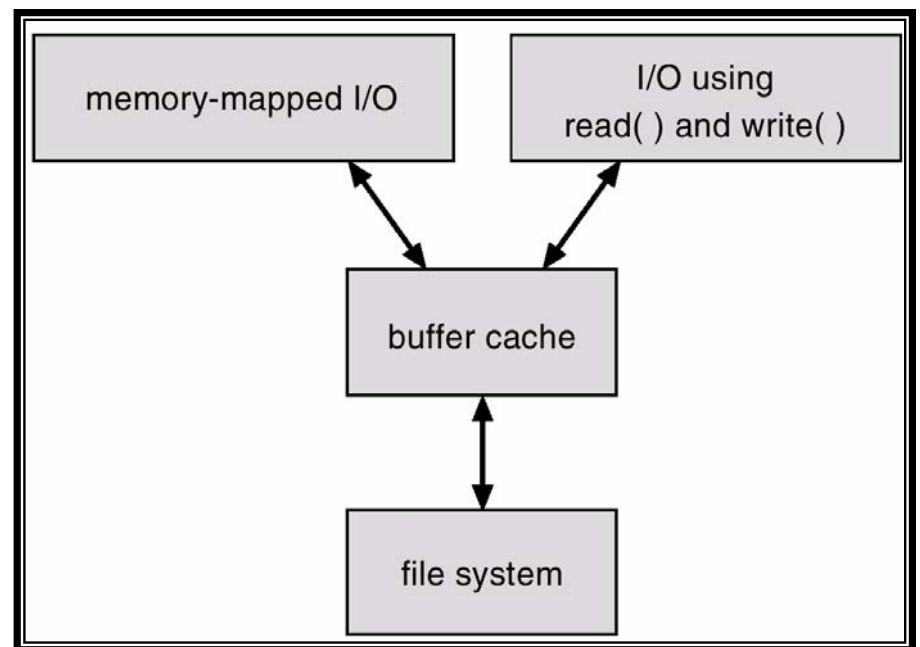
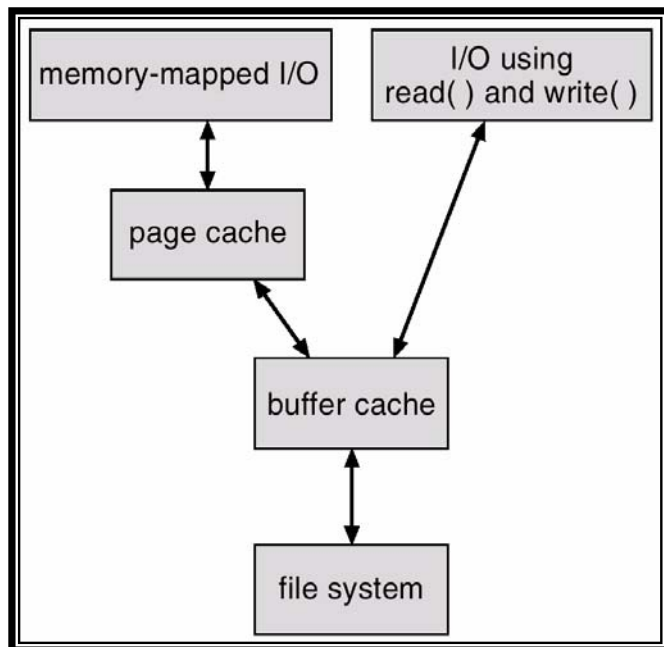


Performance (2)

- Buffer cache
 - Applications exhibit significant locality for reading and writing files.
 - Idea: cache file blocks in memory to capture locality in [buffer cache](#) (or buffer cache).
 - Cache is system wide, used and shared by all processes.
 - Reading from the cache makes a disk perform like memory.
 - Even a 4MB cache can be very effective.
 - Issues
 - The buffer cache competes with VM.
 - Live VM, it has limited size.
 - Need replacement algorithms again.
(References are relatively infrequent, so it is feasible to keep all the blocks in exact LRU order)

Performance (3)

- Unified buffer cache
 - Page caches are used to cache virtual memory pages and `mmap()`'ed file data.
 - Unified buffer cache avoids double caching and the possibility of inconsistencies between page caches and buffer caches.



Performance (4)

- Caching writes
 - Synchronous writes are very slow.
 - (1) write-behind (or asynchronous writes)
 - Maintain a queue of uncommitted blocks.
 - Periodically flush the queue to disk.
 - Unreliable: metadata requires synchronous writes. (with small files, most writes are to metadata)
 - (2) Battery backed-up RAM (NVRAM)
 - Maintain a queue in NVRAM
 - Expensive
 - (3) Log-structured file system
 - Always write next block after last block written
 - Complicated

Performance (5)

- Read ahead
 - File system predicts that the process will request next block.
 - File system goes ahead and requests it from the disk.
 - This can happen while the process is computing on previous block, overlapping I/O with execution.
 - When the process requests block, it will be in cache.
 - Compliments the disk cache, which also is doing read ahead.
 - Very effective for sequentially accessed files.
 - File systems try to prevent blocks from being scattered across the disk during allocation or by restructuring periodically.

FFS (1)

- Fast file system (FFS)
 - The original Unix file system (70's) was very simple and straightforwardly implemented:
 - Easy to implement and understand.
 - But very poor utilization of disk bandwidth (lots of seeking).
 - BSD Unix folks redesigned file system called FFS.
 - McKusick, Joy, Fabry, and Leffler (mid 80's)
 - Now it is the file system from which all other UNIX file systems have been compared.
 - The basic idea is aware of disk structure.
 - Place related things on nearby cylinders to reduce seeks.
 - Improved disk utilization, decreased response time.

FFS (2)

- Data and i-node placement
 - Original Unix FS had two major problems:
 - (1) Data blocks are allocated randomly in aging file systems.
 - Blocks for the same file allocated sequentially when FS is new.
 - As FS “ages” and fills, need to allocate blocks freed up when other files are deleted.
 - Problem: Deleted files essentially randomly placed.
 - So, blocks for new files become scattered across the disk.
 - (2) i-nodes are allocated far from blocks.
 - All i-nodes at the beginning of disk, far from data.
 - Traversing file name paths, manipulating files and directories require going back and forth from i-nodes to data blocks.
 - Both of these problems generate many long seeks!

FFS (3)

- Cylinder groups
 - BSD FFS addressed these problems using the notion of a cylinder group.
 - Disk partitioned into groups of cylinders.
 - Data blocks from a file all placed in the same cylinder group.
 - Files in same directory placed in the same cylinder group.
 - i-nodes for files allocated in the same cylinder group as file's data blocks.

FFS (4)

- Free space reserve
 - The disk must have free space scattered across all cylinders.
 - If the number of free blocks falls to zero, the file system throughput tends to be cut in half, because of the inability to localize blocks in a file.
 - A parameter, called free space reserve, gives the minimum acceptable percentage of file system blocks that should be free.
 - If the number of free blocks drops below this level, only the system administrator can continue to allocate blocks.
 - Normally 10%; this is why df may report > 100%.

FFS (5)

- Fragments
 - Small blocks (1KB) caused two problems:
 - low bandwidth utilization
 - small max file size (function of block size)
 - FFS fixes by using a larger block (4KB)
 - Allows for very large files .
(1MB only uses 2 level indirect)
 - But introduces internal fragmentation: there are many small files (i.e., < 4KB)
 - FFS introduces “fragments” to fix internal fragmentation.
 - allows the division of a block into one or more fragments (1K pieces of a block).

FFS (6)

- Media failures
 - Replicate master block (superblock)
- File system parameterization
 - Parameterize according to disk and CPU characteristics.
 - Maximum blocks per file in a cylinder group
 - Minimum percentage of free space
 - Sectors per track
 - Rotational delay between contiguous blocks
 - Tracks per cylinder, etc.
 - Skip according to rotational rate and CPU latency.

Ext2 FS (1)

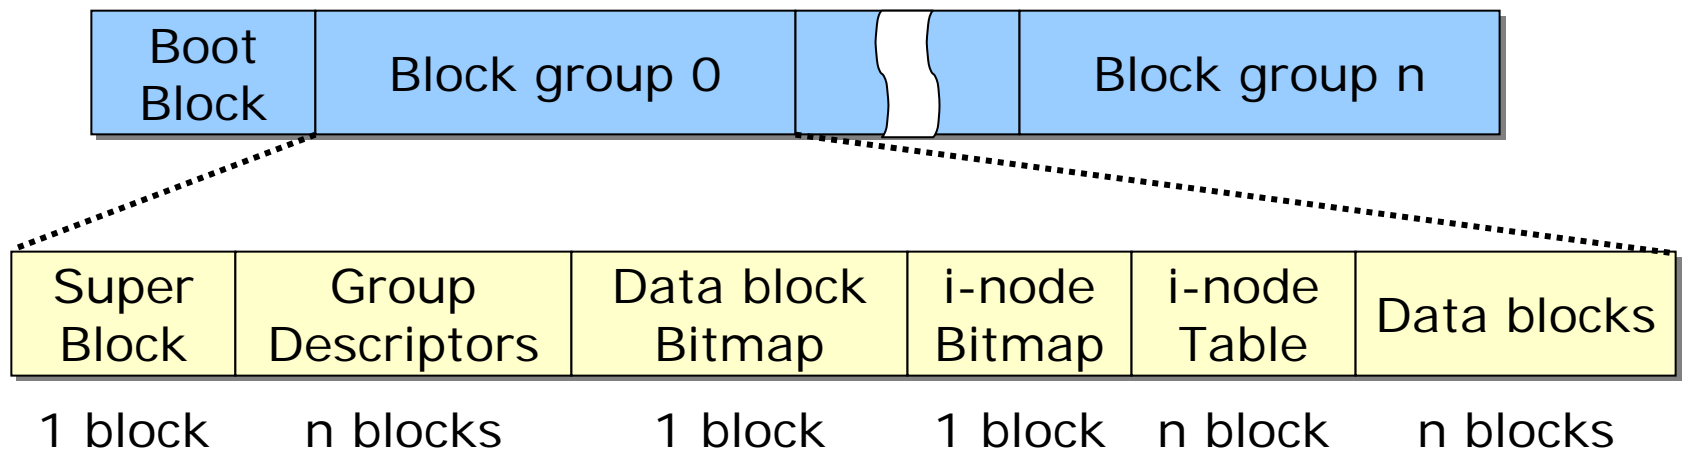
- History
 - Evolved from Minix filesystem.
 - Block addresses are stored in 16bit integers – maximal file system size is restricted to 64MB.
 - Directories contain fixed-size entries and the maximal file name was 14 characters.
 - Virtual File System (VFS) is added.
 - Extended Filesystem (Ext FS), 1992.
 - Added to Linux 0.96c
 - Maximum file system size was 2GB, and the maximal file name size was 255 characters.
 - Second Extended Filesystem (Ext2 FS), 1994.
 - Evolved to Ext3 File system (with journaling)

Ext2 FS (2)

- Ext2 Features
 - Configurable block sizes (from 1KB to 4KB)
 - depending on the expected average file size.
 - Configurable number of i-nodes
 - depending on the expected number of files
 - Partitions disk blocks into groups.
 - lower average disk seek time
 - Preallocates disk data blocks to regular files.
 - reduces file fragmentation
 - Fast symbolic links
 - If the pathname of the symbolic link has 60 bytes or less, it is stored in the i-node.
 - Automatic consistency check at boot time.

Ext2 FS (3)

- Disk layout
 - Boot block
 - reserved for the partition boot sector
 - Block group
 - Similar to the cylinder group in FFS.
 - All the block groups have the same size and are stored sequentially.



Ext2 FS (4)

- Block group
 - Superblock: stores file system metadata
 - Total number of i-nodes,
 - File system size in blocks
 - Free blocks / i-nodes counter
 - Number of blocks / i-nodes per group
 - Block size, etc.
 - Group descriptor
 - Number of free blocks / i-nodes / directories in the group
 - Block number of block / i-node bitmap, etc.
 - Both the superblock and the group descriptors are duplicated in each block group.
 - Only those in block group 0 are used by the kernel.
 - fsck copies them into all other block groups.
 - When data corruption occurs, fsck uses old copies to bring the file system back to a consistent state.

Ext2 FS (5)

- Block group size
 - The block bitmap must be stored in a single block.
 - In each block group, there can be at most $8 \times b$ blocks, where b is the block size in bytes.
 - The smaller the block size, the larger the number of block groups.
 - Example: 8GB Ext2 partition with 4KB block size
 - Each 4KB block bitmap describes 32K data blocks
 $= 32K \times 4KB = 128MB$
 - At most 64 block groups are needed.

Ext2 FS (6)

- Directory structure

	inode	record length			name			
0	21	12	1	2	.	\0	\0	\0
12	22	12	2	2	.	.	\0	\0
24	53	16	5	2	h	o	m	e
40	67	28	3	2	u	s	r	\0
62	0	16	7	1	o	l	d	f
68	34	12	3	2	b	i	n	\0

<deleted file>

name length file type